

MRP is a tuple (S, P, R, γ) , where S is the set of states P is 2-dimensional transition matrix R is a reward function defined as $R = E[r_{t+1} | S_t = s]$, that is collected upon leaving a the state s , i.e. at time step $t+1$ γ is the discount factor

Return

- Total discounted reward at from time step t

$$R_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

- γ close to 0 leads to short-sighted evaluation (i.e. don't look into the future too much)
- γ close to 1 leads to far-sighted evaluation This return is specific for some sample of states. It does not take into account all possible paths that can be taken from a given state

Value

The aggregated return for a given state can be thought of as its 'value'

$$v(s) = E[R_t | S_t = s]$$

Bellman Equation for MRP

$$\begin{aligned} v(s) &= E[R_t | S_t = s] \quad \&= E[r_{t+1} + \gamma v(s_{t+1}) | S_t = s] \\ &= E[r_{t+1} + \gamma (R_{t+1} + \gamma v(s_{t+1})) | S_t = s] \end{aligned}$$

The transition from (1) to (2) uses the Tower rule and in essence the conditional in the inner expected value will be averaged out/canceled out so it's value does not matter. Therefore it is more useful to use the value of the next state $v(s_{t+1})$

See proof: <https://chat.openai.com/share/77adb96c-1d17-4fca-8d88-91792567c997>

Different Notation

Sum notation (written out)

$$v(s) = R_s + \gamma \sum_{s'} P_{ss'} v(s')$$

Vector form

$$v = R + \gamma P v$$

This can be analytically solved:
$$\begin{aligned} v &= R + \gamma P v \\ v - \gamma P v &= R \\ (\mathbb{I} - \gamma P) v &= R \\ v &= (\mathbb{I} - \gamma P)^{-1} R \end{aligned}$$

Matrix inversion however has a complexity of $O(n^3)$ and therefore works only for small MRPs

Policy (π)

MRP helps us observe the state space and figure out the value of each state. In order to actually influence the environment and take actions a policy π is introduced.

A policy could either be deterministic and non-deterministic.